

Seeing more with less scanning

What happens when an electron microscope visits only 5–30% of the pixels, and which ways of filling in the rest actually work.

A plain-language summary of the full study in the [sparse-scan-SEM-reconstruction](#) repository. Every number here comes from its committed results. 100 test images, 6 methods, 30 test conditions per image.

The short version

1. **The trained neural network (U-Net) is the clear winner.** It is the best method in every one of the 30 test conditions, about 2 dB ahead of the best traditional method on average, better on 99% of individual images, and it takes 0.03 seconds per image.
2. **How you skip pixels matters as much as how many.** Scanning a subset of full lines (partial raster) gives the best picture for the time spent, by a wide margin. Scattering pixels at random wastes most of the time moving the beam.
3. **The fancier diffusion model comes second.** It is 0.4 dB behind the U-Net and 160× slower, at least on the standard accuracy scores.
4. **A network must be trained on the scan patterns it will see.** One trained only on random pixels falls apart on line-based scans.

The problem

A scanning electron microscope (SEM) builds an image one pixel at a time. It points a beam at a spot, counts the electrons that bounce back, and moves on. A 1024×768 image means about 800,000 of those visits. Two things make skipping most of them attractive:

- **Speed.** Fewer visits means a faster scan, which matters for large areas, moving processes and busy instruments.
- **Dose.** The beam damages delicate specimens such as biological tissue and some polymers. Fewer visits means less damage.


The catch is that the skipped pixels then have to be filled in by software. This study asks which software does that best, and which way of choosing pixels gives the software the best chance.

How the experiment works

Everything is simulated on real SEM images, so the true image is always known and every method can be scored exactly. The images come from the public NFFA-Europe "100% SEM" collection: 21,169 micrographs in 10 categories, from biological samples to MEMS devices.

1. Pick which pixels to measure

Three sampling patterns are compared. Each one measures the same share of pixels, called the *sampling fraction* (5%, 10%, 15%, 20% or 30%).



The three patterns at 10% of pixels. **Uniform**: random scattered pixels. **Partial raster**: a subset of complete scan lines. **Line-hop**: every line is visited, but the beam only stops on random short segments and hops forward over the gaps. The time above each panel comes from the scan-time model below.

2. Charge for the time the beam really takes

A beam can't teleport. Every jump to a new spot needs the deflection coils to settle, and every new line needs a flyback to the start. The simulation charges 1 μ s per measured pixel, 5 μ s per jump and 50 μ s per line. That changes the picture completely. At 10% of pixels, random scattering costs 0.58 $\times$ the time of a full scan, because almost every pixel needs a jump. Line-hop costs 0.20 $\times$, and partial raster only 0.10 $\times$. At 20% random pixels, the time saving is gone entirely.

3. Add realistic noise

Electron counts are random, like raindrops in a bucket: expect 20 and you might get 16 or 25. This "shot noise" is simulated faithfully. Two scenarios are tested:

- **Fixed dwell**: every measured pixel gets the same exposure, 20 electrons on average for a bright spot. This asks: how much faster can I scan at the same per-pixel quality?
- **Fixed total dose**: the specimen gets the same total exposure as a quick full scan (2 electrons per pixel), shared among whichever pixels are measured. Measuring fewer pixels means each one gets more electrons and is less noisy. This asks: for a fragile specimen with a fixed damage budget, is it better to measure a few pixels well or many pixels roughly?

4. Fill in the gaps and score the result

Each method reconstructs the full image, and it is scored against the true image. All methods see exactly the same measured pixels and the same noise, so the comparison is fair. The test set is 100 image crops of 512×512 pixels. With 3 patterns × 5 fractions × 2 scenarios, each method produces 3,000 reconstructions.

Reading the scores. The main score is **PSNR**, in decibels (dB). Higher is better, and it measures how close the reconstruction is to the truth, pixel by pixel. It's a log scale:

- +1 dB means about 21% less error;
- +2 dB means about 37% less;
- +3 dB means half the error.

SSIM, from 0 to 1, measures how well the structure is preserved, such as edges and textures.

The contenders

Biharmonic	Draws the smoothest possible surface through the measured pixels. Simple and exact, but it passes the measurement noise straight through.
TV-L2	A classic "total variation" method. It looks for an image that agrees with the measurements but has as few sharp changes as possible, which removes noise but also flattens fine texture. Its strength setting was tuned for each condition on separate validation images.
TV-Poisson	The same idea, with a noise model that matches electron counting exactly. In principle it should win when counts are low.
U-Net	A neural network (7.8M parameters) trained on thousands of SEM images to turn "measured pixels plus a map of where they are" into a full image. One network handles every pattern, fraction and noise level.
U-Net, uniform-only	The same network trained only on random-pixel masks. It's a control, to test whether training on the right patterns matters.
Diffusion	A generative model (7.9M parameters) of the kind behind modern AI image generators. It starts from pure noise and gradually "develops" an image consistent with the measurements. Run it several times and you get different plausible images; the score here uses the average of 4.


What we found

1. The neural network wins everywhere, and it's the fastest.

Averaged over all sampling fractions with fixed dwell, the U-Net scores about 2 dB above TV-L2, the best traditional method. That's roughly 37% less error. The gap holds on 99% of individual test images and in every one of the 30 conditions. In SSIM, the structure score, the U-Net also leads, at 0.58 against 0.48 for TV-L2.

Method	Uniform	Partial raster	Line-hop	Time per image
Biharmonic	18.03	19.66	16.75	~12 s (CPU)
TV-L2	24.05	23.53	23.21	~12 s (CPU)
TV-Poisson	23.70	23.25	22.90	~12 s (CPU)
U-Net, uniform-only	25.90	22.79	23.53	0.03 s (GPU)
U-Net	26.11	25.48	25.22	0.03 s (GPU)
Diffusion	25.79	25.08	24.83	4.8 s (GPU)

PSNR in dB, averaged over the five sampling fractions, fixed dwell. Time is for one 512×512 image.




One test image from 10% line-hop measurements. Biharmonic is full of noise. Both TV methods wash out the texture and leave streaks. The uniform-only U-Net invents vertical streaks, because it never saw line-hop masks in training. The U-Net and diffusion both recover the fibres.

2. For the time spent, scan whole lines.

This is the most practical result. Give each method a time budget and let it pick whichever pattern and fraction does best within it. At every budget up to a full scan's time, the answer is partial raster, for every method except the uniform-only control. With the U-Net, 30% of lines takes 0.30× the time of a full scan and reaches 27.1 dB. Random scattering at just 5% already takes longer (0.34×) and reaches only 24.6 dB.

Past about 0.3× the time, extra scanning buys almost nothing. The only configuration that edges past 30% of lines is random scattering at 30%, and it takes 1.31× as long as scanning every pixel.



The best score each method can reach within a scan-time budget (fixed dwell). Each step is a better configuration becoming affordable, and almost all of them are partial raster. The horizontal axis is time relative to a full scan.

3. The diffusion model is good, but not better on these scores.

Diffusion is a solid second: 1.6 dB ahead of TV-L2, but 0.37–0.41 dB behind the U-Net on 98–99% of images. It's also 160× slower, because it runs its network 100 times per sample, 4 samples per image.

Why doesn't the fancier model win? PSNR rewards the *average* of all plausible answers, and that's exactly what the U-Net is trained to predict. A diffusion model is built to produce individual, realistic-looking answers, and averaging 4 of them only gets close to the true average. Its real strengths, realistic texture and a map of its own uncertainty, don't show up in PSNR. A follow-up experiment measuring exactly those is running now.


One design detail mattered a lot. Forcing the diffusion model to copy the measured pixels exactly made it about 5 dB *worse*, because those pixels are noisy. Letting the network treat them as hints worked best.

4. Train on the patterns you'll use.

The control network, trained only on random pixels, nearly matches the full U-Net on random pixels (0.2 dB behind). On line-based scans, though, it loses 1.7–3.1 dB, and on partial raster it drops below even the traditional TV-L2. A network only knows the kinds of gaps it was shown.

5. With a fixed damage budget, only noise-aware methods benefit from more pixels.

When the total dose is fixed, measuring more pixels means each one is noisier. The U-Net keeps improving anyway, from 24.3 dB at 5% to 25.8 dB at 30%, because it removes noise as it fills in. Biharmonic, which copies noise exactly, gets steadily *worse*, from 19.1 to 14.6 dB. For the learned models the gains level off beyond about 15%.




Fixed total dose. Each panel is one sampling pattern, and each line one method. Shaded bands show the spread across test images: some images are simply harder than others. Head-to-head comparisons on the same image are far more consistent than the bands suggest.

6. The exact noise model didn't help the traditional method.

TV-Poisson, which models electron counting exactly, came out 0.3–0.4 dB *behind* the simpler TV-L2, even at the lowest counts. We checked this wasn't a tuning mistake by widening its settings range; one case improved, but the conclusion held. Likely reasons: at these counts the simple approximation is already good enough, and the "true" images used for scoring are themselves slightly noisy.

7. Some specimens are much harder than others, but the ranking never changes.

Fine, dense textures such as porous sponges and coated films are hardest, at about 21 dB for the U-Net. Isolated objects on smooth backgrounds, such as probe tips and particles, are easiest, at about 30 dB. In all 10 categories the order is the same: U-Net, then diffusion, then the traditional methods. The U-Net's advantage is largest on fibres (+3.2 dB over TV-L2) and smallest on biological samples (+1.2 dB).



Average score per specimen category, 10 test images each. The U-Net (orange diamond) is best in every row.

What this means in practice

- **For a faster scan**, skip whole lines rather than scattered pixels, and use a trained network to fill them in. Scanning 20–30% of lines gave the best results we measured, in a fifth to a third of the time of a full scan.
- **For a fragile specimen**, spreading a fixed dose over more pixels is fine, even helpful, provided the reconstruction removes noise. Simple interpolation is the wrong tool there.
- **If you train a network**, train it on the scan patterns your microscope will actually use.

Limits of this study

- **It's a simulation.** Real deflection coils overshoot and wobble after jumps, the stage drifts, and some specimens build up electric charge. None of that is modelled yet. A follow-up is testing how coil position errors would change the results.
- **The "truth" isn't perfect.** The reference images are themselves ordinary, slightly noisy, compressed micrographs. Scores measure agreement with them, and the networks partly learn their compression artefacts.

- **Accuracy scores aren't everything.** PSNR and SSIM favour smooth, averaged answers. A method that produces sharper, more realistic texture can score lower and still look better to a person.

What's in progress

- **What diffusion is actually good for.** Does its uncertainty map point to where it's wrong? Do single samples keep more real texture than the U-Net? Is averaging 4 samples enough? This is running on the GPU now.
- **Imperfect scan coils.** When the beam lands slightly off target after each jump, do random-pixel scans suffer most, as expected, since they jump before almost every pixel? This would strengthen the case for line-based scanning.

The trained models, all code and every result table are public in the repository, under the Apache 2.0 license.

Glossary

Sampling fraction	The share of pixels actually measured, 5–30% here.
Raster scan	The normal way an SEM scans: line by line, left to right, like reading a page.
Shot noise	Randomness in how many electrons arrive, even under steady conditions. More electrons per pixel means relatively less noise.
Dwell / dose	How long the beam stays on a pixel, and so how many electrons hit it. Dose is the total exposure the specimen receives.
PSNR	Peak signal-to-noise ratio, in dB. A logarithmic score of closeness to the truth; +3 dB means half the error.
SSIM	Structural similarity, from 0 to 1. How well edges, contrast and texture are preserved.
U-Net	A common neural-network design for image-to-image tasks, shaped like a "U": it shrinks the image to understand context, then expands it back to full detail.
Diffusion model	A generative network that creates images by repeatedly removing noise, guided here by the measured pixels.
Ablation	A control experiment that removes one ingredient, here training on line patterns, to see how much it matters.